

# Week 6 and 7 Recap

Pranav Rao

October 19, 2023

## 1 Preface

This document will cover week 6 and 7 of the STA237H1 curriculum. Some of it will be copy-paste from my Week 3 to 5 recap, but this document will omit information about discrete distributions and focus more on continuous distributions.

## 2 What is a continuous random variable?

A continuous random variable can be defined as “a random variable that can take on an infinite number of possible values”.

## 3 PDFs and CDFs

For any continuous statistical distribution, there are two main functions that can be used to describe it.

- The first function we can use to describe a continuous distribution is called a **probability density function** (PDF for short, not to be confused with the filetype). This function is often denoted  $f_X(k)$ .
  - While a PMF tells you the probability of a random variable  $X$  being a certain value  $k$ , the PDF  $f_X(k)$  represents the *height of the graph of  $X$ 's values at  $k$* , meaning you have to take the area *under* the graph to get anything useful.
    - \* Note: the area under a graph at a *specific value* is 0, which means it always holds that  $f_X(k) = 0$  for a continuous  $X$  and a specific  $k$ .
    - \* This is why, more often than not, with continuous distributions, we deal with *integrations* of the PDF over ranges (e.g. the CDF) rather than the PDF itself.
    - \* As a consequence of  $f_X(k) = 0$ ,  $P(X < k) = P(X \leq k)$
  - The PDF,  $f_X$ , for a continuous variable  $X$  satisfies the following conditions:

- \*  $f_X(x) > 0$
- \*  $f_X$  is piecewise continuous
- \*  $\int_{-\infty}^{\infty} f_X(x)dx = 1$
- The second function used to describe a continuous distribution is called a **cumulative distribution function** (CDF). The CDF is a function such that, when evaluated at some  $x$ , it gives the cumulative probability that the random variable  $X$  will take a value less than or equal to  $X$ . For a continuous random variable  $X$  with PDF  $f_X$ , the CDF  $F_X(x)$  is calculated as such:

$$F_X(x) = P(X \leq x) = \int_{-\infty}^x f_X(y)dy$$

## 4 Expected Values, Variances, and Standard Deviation of Continuous Random Variables

This section will cover how to get the expected value, variance, and standard deviation of continuous random variable in the general case. See Week 3 to 5 Recap for a more detailed definition of these terms.

### 4.1 Expected Values of Continuous Random Variables

The expected value of a continuous random variable  $X$ , given that its PDF is  $f_X(x)$  for some value  $x$ , can be calculated as such

$$E[X] = \mu = \int_{-\infty}^{\infty} x \cdot f_X(x)dx$$

## 5 Variance and Standard Distribution

Variance is calculated in the same way for both types of random variables. As a recap:

$$Var(X) = \sigma^2 = E[(X - \mu)^2] = E[X^2] - (E[X])^2$$

Also, standard deviation, represented by  $\sigma$  is still the square root of the variance:

$$\sigma = \sqrt{\sigma^2} = \sqrt{E[(X - \mu)^2]} = \sqrt{E[X^2] - (E[X])^2}$$

Refer to the Week 3 to 5 Recap for more information.

## 6 Moment Generating Functions

Refer to the Week 3 to 5 Recap for context on what a moment-generating function is. For a continuous random variable  $X$  with PDF  $f_X$ , the moment generating function is defined as:

$$M_X(t) = E(e^{tX}) = \int_{-\infty}^{\infty} e^{tx} f_X(x) dx$$

## 7 Quantiles and Percentiles

### 7.1 Quantiles

- A **quantile** is a cut-point dividing the range of a probability distribution or dividing observations in a sample into continuous intervals with equal probabilities.
- **$q$ -quantiles** are values that partition a *finite* set into  $q$  subsets of (nearly) equal sizes; there are  $q - 1$  partition points to form  $q$  partitions.
- Some  $q$ -quantiles have special names:
  - The singular 2-quantile is called the **median**.
  - The two 3-quantiles are called tertiles
  - The three 4-quantiles are called quantiles
  - So on and so-forth for 5, 6, 7...-quantiles
- Quantiles are expressed in the unit of measurement of the input scores, not in percentages.

### 7.2 Percentiles

- In essence, a **percentiles** are 100-quantiles.
  - The 25th percentile is the *first quartile* (see above)
  - The 50th percentile is the *median*
  - The 75th percentile is the *third quartile*
- Independent of a quantile, a the  **$k$ -th percentile** (also known as the percentile score or the centile) is defined as a score (at or) below which a given percentage falls (depends on whether or not you are looking at *inclusive* or *exclusive* percentiles).
  - For example: suppose there was a test where the 90th percentile was a 75%. This means that 90% of the scores on the test were (at or) below 75%.

- Like quantiles, are expressed in the unit of measurement of the input scores, not in percentages.
- You can calculate percentiles on a distribution using the `q` variant of the distribution function (see “R Distribution Functions” in the week 3 to 5 recap).

### 7.3 The Quantile Function

- The **quantile function** of a probability distribution (also known as the **percentile function**) is a function that, given an input probability value, outputs the probability that the value of a random variable with that probability distribution will be less than or equal to that input probability value.
- This is the same as a distribution’s **inverse cumulative distribution function**, which intuitively make sense because it does the opposite of the CDF.
- See the “R Distribution Functions” from the Week 3 to 5 recap.

## 8 Common Continuous Distributions

Note: the examples in this section are examples with the PDF; however, in practice, for continuous distributions, we generally work with ranges and practically using the CDF is much more common.

### 8.1 Continuous Uniform Distribution

- *Use case*: where there is an arbitrary outcome that lies within certain bounds.
- *Examples*: you’re waiting for a bus that comes once hourly, but you don’t know when it came last (where  $a = 0$ ,  $b = 1$  measured in hours)
- *Notation*:  $Unif(a, b)$ , where  $a$  is the first lower bound and  $b$  is the upper bound
- *PDF*:  $f_X(k) = \begin{cases} \frac{1}{b-a} & \text{if } k \in [a, b] \\ 0 & \text{otherwise} \end{cases}$ , where  $a$  is the first lower bound and  $b$  is the upper bound
- *Mean*:  $\frac{a+b}{2}$ , where  $a$  is the lower bound and  $b$  is the upper bound
- *Variance*:  $\frac{1}{12}(b-a)^2$ , where  $a$  is the lower bound and  $b$  is the upper bound
- *Associated R functions*: `dunif`, `punif`, `qunif`, `runif`

## 8.2 Exponential Distribution

Note: this distribution is sometimes also called the negative exponential distribution.

- *Use case*: where you have a situation where events occur continuously and independently at a constant average rate  $\lambda$ , and you care about the wait time until a first event.
- *Examples*: (not sure, hard to find examples for PDF) Suppose you wanted to know the probability of taking between 3 and 4 seconds to burn out, given it generally fails after 5 seconds (where  $\lambda = 5$ ).
- *Notation*:  $\text{Exp}(\lambda)$ , where  $\lambda$  is the mean or the rate
- *PDF*:  $f_X(k) = \lambda e^{-\lambda k}$ , where  $\lambda$  is the mean or the rate
- *Mean*:  $\frac{1}{\lambda}$ , where  $\lambda$  is the the mean or the rate
- *Variance*:  $\frac{1}{\lambda^2}$ , where  $\lambda$  is the mean or the rate
- *Associated R functions*: `dexp`, `pexp`, `qexp`, `rexp`

## 8.3 Normal Distribution

- *Use case*: as input increases, if the output is low, high, and then low again (most people around the mean)
- *Examples*: Suppose you given are the average height of the population is 5 feet and the average standard deviation of heights is 0.5, and you want to find the probability that someone has a height between 5.0 and 5.1 feet (where  $\mu = 5, \sigma = 0.5$ )
- *Notation*:  $N(\mu, \sigma^2)$ , where  $\mu$  is the mean and  $\sigma^2$  is the variance
- *PDF*:  $f_X(k) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}(\frac{k-\mu}{\sigma})^2}$
- *Mean*:  $\mu$
- *Variance*:  $\sigma^2$
- *Associated R functions*: `dnorm`, `pnorm`, `qnorm`, `rnorm`

## 8.4 Gamma Distribution

- *Use case:* where you have a situation where events occur continuously and independently at a constant average rate  $\lambda$ , and you care about the wait time until the  $k$ 'th event.
- *Examples:* want to find the probability a machine takes between 3 and 4 units of time to produce a specified number of items given we are dealing with 5 units of time overall and it produces 2 items per unit of time.
- *Notation:*  $\text{Gam}(\alpha, \lambda)$ , where  $\alpha$  is shape parameter (number of events you are modelling) and  $\lambda$  is the rate parameter (the rate at which constantly-spaced events occur)
- *PDF:*  $f_X(k) = \frac{\lambda(\lambda k)^{\alpha-1} e^{-\lambda k}}{\Gamma(\alpha)}$ , where  $\Gamma(\alpha) = (\alpha - 1)!$  for any  $\alpha \in \mathbb{N}, \alpha > 0$ .
- *Mean:*  $\frac{\alpha}{\lambda}$ , where  $\alpha$  is shape parameter (number of events you are modelling) and  $\lambda$  is the rate parameter (the rate at which constantly-spaced events occur)
- *Variance:*  $\frac{\alpha}{\lambda^2}$ , where  $\alpha$  is shape parameter (number of events you are modelling) and  $\lambda$  is the rate parameter (the rate at which constantly-spaced events occur)
- *Associated R functions:* `dgamma`, `pgamma`, `qgamma`, `rgamma`

## 8.5 Cauchy Distribution

It is worth noting this distribution is in the textbook. However, it does not at all follow the rules of a regular distribution because it has no moments. Furthermore, examples are not easy to find and it is useless for the purposes of problems in this course, so I have chosen to omit it as of now.

## 8.6 Pareto Distribution

Also known as the “80-20 rule”. Note: this version is from the textbook, but the department wants us to use a 2-parameter Pareto distribution so I will update this at one point.

- *Use case:* where you want to understand the probability of a specific event occurring at a particular threshold or level in a system. The distribution is right-skewed and has a heavy long tail. Example: it is useful for modelling catastrophic events where a claim might have a very large value.
- *Examples:* want to find the probability that a person has between a \$1000 and \$2000 income, given that people’s incomes follow a Pareto distribution with a rate parameter  $\alpha = 3$  and a location parameter  $\beta = 1$ . (*Note: you pretty much need to be given that a situation follows a Pareto distribution with a given shape parameter to actually use the distribution, it’s not very intuitive.*)

- *Notation:*  $Par(\alpha, \beta)$ , where  $\alpha$  is shape/rate parameter (slope of the distribution/how quickly it drops off; low  $\alpha$  implies low drop off implies that more extreme outcomes are more likely).  $\beta$  is the location parameter. Note the Pareto distribution in the textbook is a  $Par(\alpha, 1)$  distribution.
- *PDF:*  $f_X(k) = \frac{\alpha\beta^\alpha}{k^{\alpha+1}}$  for  $k \geq \beta$ , where  $\alpha$  is the shape parameter and  $\beta$  is the location parameter
- *Mean:*  $\begin{cases} \frac{\alpha\beta}{\alpha-1} & \text{if } \alpha > 1 \\ \infty & \text{if } 0 < \alpha \leq 1 \end{cases}$ , where  $\alpha$  is the shape parameter
- *Variance:*  $\begin{cases} \frac{\alpha\beta^2}{(\alpha-1)^2(\alpha-2)} & \text{if } \alpha > 2 \\ \infty & \text{if } 0 < \alpha \leq 1 \end{cases}$
- *Associated R functions:* `dpareto` `ppareto`, `qpareto`, `rpareto` (not in standard library, install and import `EnvStats`)

## 8.7 Beta Distribution

- *Use case:* when you want the probability that something is  $x\%$  successful
- *Example:* What is the probability a website's conversion, which follows a beta distribution with parameters  $\alpha = 5, \beta = 10$  is between 60% and 70%? (*Note: similarly to the Pareto distribution, you pretty much have to be told the variable follows a Beta distribution*).
- *Notation:*  $B(\alpha, \beta)$ , where  $\alpha$  and  $\beta$  are shape parameters.
- *PDF:*  $f_X(k) = \begin{cases} \frac{1}{B(\alpha, \beta)} x^{\alpha-1} (1-y)^{\beta-1} & 0 \leq y \leq 1 \\ 0 & \text{otherwise} \end{cases}$  where  $B(\alpha, \beta) = \int_0^1 x^{\alpha-1} (1-x)^{\beta-1} dy = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha+\beta)}$ , where  $\alpha, \beta$  are shape parameters and  $\Gamma$  is the gamma function.
- *Mean:*  $\frac{\alpha}{\alpha+\beta}$ , where  $\alpha, \beta$  are the shape parameters
- *Variance:*  $\frac{\alpha\beta}{(\alpha+\beta+1)(\alpha+\beta)^2}$
- *Associated R functions:* `dbeta` `pbeta`, `qbeta`, `rbeta`

## 9 Histograms and Probability Plots

There are different techniques to visualize data and understand which distribution might best fit the data. This section covers two of them.

## 9.1 Histograms

- For large sample sizes, typically **histograms** are used to determine which distribution fits the data
  - A **histogram** is a commonly used graph to show *frequency distributions* (i.e. how often each value in a dataset shows up in that dataset)
  - Example histogram:




Figure 1: An example histogram

## 9.2 Probability Plots

- For smaller sample sizes, histograms are not as effective, so a graphical method we can use instead is probability plots.
  - A **probability plot** is used to confirm a hypothesis about whether or not some data follows a certain distribution.
- A worked example of how to use a probability plot to confirm if some data has a normal distribution:
  1. Suppose we have some data, which can be found in the table below:



---

Observed Length of Fish

---

100  
98  
101  
93  
123  
112  
85  
76  
119  
111

2. Sort the data in ascending order:

---

Observed Length of Fish

---

76  
85  
93  
98  
100  
101  
111  
112  
119  
123

3. Number the data from index 1 to 10:

Number (i)	Observed Length of Fish
1	76
2	85
3	93
4	98
5	100
6	101
7	111
8	112
9	119
10	123

4. For each index, calculate the “Expected Cumulative Probability”, which follows the formula  $\frac{i-0.5}{n}$ , where  $i$  is the index in question and  $n$  is the total data points:

Number (i)	Observed Length of Fish	Expected Cumulative Probability
1	76	0.05
2	85	0.15
3	93	0.25
4	98	0.35
5	100	0.45
6	101	0.55
7	111	0.65
8	112	0.75
9	119	0.85
10	123	0.95

5. Determine the  $z$ -scores (or  $z$ -value) for each expected cumulative probability from the standard normal probability distribution (which is a normal probability distribution with  $\mu = 0$  and  $\sigma = 1$ ). That is, you need to find which  $z$ -score will give you the expected cumulative probability in question. Some notes about  $z$ -scores before the table containing the  $z$ -scores for cumulative probabilities in this example:

- A  $z$ -score/value is a statistical measurement that describes the value's relationship to the mean of a group of values.  $z$ -scores are measured in terms of “standard deviations from the mean.” For example, if a value of the random variable  $X = x$  has a  $z$ -score of 1, it means that the value of  $X = x$  is 1 standard deviation to the right of the mean.
- (*Not useful for this example but useful to know in general*) Given a value  $X = x$ , a  $z$ -score for a normal distribution can be calculated with the following formula (where  $\mu$  is the mean, and  $\sigma$  is the standard deviation of the distribution):

$$z = \frac{x - \mu}{\sigma}$$

- (*Somewhat useful for this example*) Given a  $z$ -value, if you want to find the area to the left of that  $z$  value (essentially area to the left of  $z$  standard deviations from the mean), you can use a **Z lookup table** (which would have to be provided to us, see this link) for information.
- (*Directly useful for this example*): in this case, we want to find the opposite of the above point; we want to find the  $z$  *given* the area to the left of the  $z$ -score (the Expected Cumulative Probability). We can do this in one of two ways:
  - \* Use the *normal distribution's quantile function* (very complicated, would not even bother looking it up)
  - \* Use the *Z lookup table, but backwards*
  - \* (Recommended) Use software, namely the `qnorm` function from R, which does exactly what we want.

Plugging in each Expected Cumulative Probability into **qnorm**:

Number (i)	Observed Length of Fish	Expected Cumulative Probability	Z-Scores (approx)
1	76	0.05	-1.64
2	85	0.15	-1.03
3	93	0.25	-0.67
4	98	0.35	-0.39
5	100	0.45	-0.13
6	101	0.55	0.13
7	111	0.65	0.38
8	112	0.75	0.67
9	119	0.85	1.04
10	123	0.95	1.64

- Continuing on:

6. Plot the points  $(x, y)$ , where  $x$  is the Observed Length of Fish and  $y$  is the associated  $z$ -score. This plot is a **normal probability plot**. If the points fall roughly in a straight line, you can assume you have a normal distribution, and your initial hypothesis is correct.

- For my sanity and yours, I will not include examples for more complicated distributions (like exponential, etc.). They are extremely hard to find and I doubt we will be tested on them because normal probability plots are hard enough.

## 10 Simulating Realizations of Random Variables with the Uniform Distribution

Consider the following theorem:

**Theorem** (Inverse Transform Method). *Suppose  $X$  is a continuous random variable with CDF  $F$ , where  $F$  is invertible with inverse function  $F^{-1}$ . Let  $U \sim \text{Unif}(0, 1)$ . Then the distribution of  $F^{-1}(U)$  is equal to the distribution of  $X$ .*